

Name Bealu Girma Gebresilassie

Roll no: 21053273 CSE-2

In []: `!pip install transformers "datasets[s3]==2.18.0" "sagemaker>=2.190.0" "hug`

```
Requirement already satisfied: transformers in /usr/local/lib/python
3.10/dist-packages (4.38.2)
Collecting transformers
  Downloading transformers-4.39.1-py3-none-any.whl (8.8 MB)
      8.8/8.8 MB 18.6 MB/s et
a 0:00:00
Collecting datasets[s3]==2.18.0
  Downloading datasets-2.18.0-py3-none-any.whl (510 kB)
      510.5/510.5 kB 31.3 MB/
s eta 0:00:00
Collecting sagemaker>=2.190.0
  Downloading sagemaker-2.214.1-py3-none-any.whl (1.4 MB)
      1.4/1.4 MB 36.1 MB/s et
a 0:00:00
Requirement already satisfied: huggingface_hub[cli] in /usr/local/li
b/python3.10/dist-packages (0.20.3)
Collecting huggingface_hub[cli]
  Downloading huggingface_hub-0.22.1-py3-none-any.whl (388 kB)
      388.6/388.6 kB 28.4 MB/
+ . . . .
```

In []: `!huggingface-cli login --token hf_jwOUrKKKVGoOUrnBQZZQgJTPhyTHqidWd`

```
Token has not been saved to git credential helper. Pass `add_to_git_cred
ential=True` if you want to set the git credential as well.
Token is valid (permission: read).
Your token has been saved to /root/.cache/huggingface/token
Login successful
```

In []: ▶ !aws configure

Requirement already satisfied: sagemaker in /usr/local/lib/python3.10/dist-packages (2.214.1)

Requirement already satisfied: attrs<24,>=23.1.0 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (23.2.0)

Requirement already satisfied: boto3<2.0,>=1.33.3 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (1.34.72)

Requirement already satisfied: cloudpickle==2.2.1 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (2.2.1)

Requirement already satisfied: google-pasta in /usr/local/lib/python3.10/dist-packages (from sagemaker) (0.2.0)

Requirement already satisfied: numpy<2.0,>=1.9.0 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (1.25.2)

Requirement already satisfied: protobuf<5.0,>=3.12 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (3.20.3)

Requirement already satisfied: smdebug-rulesconfig==1.0.1 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (1.0.1)

Requirement already satisfied: importlib-metadata<7.0,>=1.4.0 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (6.11.0)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (24.0)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from sagemaker) (1.5.3)

Requirement already satisfied: pathos in /usr/local/lib/python3.10/dist-packages (from sagemaker) (0.3.2)

Requirement already satisfied: schema in /usr/local/lib/python3.10/dist-packages (from sagemaker) (0.7.5)

Requirement already satisfied: PyYAML~=6.0 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (6.0.1)

Requirement already satisfied: jsonschema in /usr/local/lib/python3.10/dist-packages (from sagemaker) (4.19.2)

Requirement already satisfied: platformdirs in /usr/local/lib/python3.10/dist-packages (from sagemaker) (4.2.0)

Requirement already satisfied: tblib<4,>=1.7.0 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (3.0.0)

Requirement already satisfied: urllib3<3.0.0,>=1.26.8 in /usr/local/lib/python3.10/dist-packages (from sagemaker) (2.0.7)

Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from sagemaker) (2.31.0)

Requirement already satisfied: docker in /usr/local/lib/python3.10/dist-packages (from sagemaker) (7.0.0)

Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from sagemaker) (4.66.2)

Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from sagemaker) (5.9.5)

Requirement already satisfied: botocore<1.35.0,>=1.34.72 in /usr/local/lib/python3.10/dist-packages (from boto3<2.0,>=1.33.3->sagemaker) (1.34.72)

Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in /usr/local/lib/python3.10/dist-packages (from boto3<2.0,>=1.33.3->sagemaker) (1.0.1)

Requirement already satisfied: s3transfer<0.11.0,>=0.10.0 in /usr/local/lib/python3.10/dist-packages (from boto3<2.0,>=1.33.3->sagemaker) (0.10.1)

Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.10/dist-packages (from importlib-metadata<7.0,>=1.4.0->sagemaker) (3.18.1)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->sagemaker) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10

```

0/dist-packages (from requests->sagemaker) (3.6)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->sagemaker) (2024.2.2)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from google-pasta->sagemaker) (1.16.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /usr/local/lib/python3.10/dist-packages (from jsonschema->sagemaker) (2023.12.1)
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/python3.10/dist-packages (from jsonschema->sagemaker) (0.34.0)
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.10/dist-packages (from jsonschema->sagemaker) (0.18.0)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->sagemaker) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->sagemaker) (2023.4)
Requirement already satisfied: ppft>=1.7.6.8 in /usr/local/lib/python3.10/dist-packages (from pathos->sagemaker) (1.7.6.8)
Requirement already satisfied: dill>=0.3.8 in /usr/local/lib/python3.10/dist-packages (from pathos->sagemaker) (0.3.8)
Requirement already satisfied: pox>=0.3.4 in /usr/local/lib/python3.10/dist-packages (from pathos->sagemaker) (0.3.4)
Requirement already satisfied: multiprocessing>=0.70.16 in /usr/local/lib/python3.10/dist-packages (from pathos->sagemaker) (0.70.16)
Requirement already satisfied: contextlib2>=0.5.5 in /usr/local/lib/python3.10/dist-packages (from schema->sagemaker) (21.6.0)

```

```

In [ ]: ▶ import sagemaker
import boto3
sess = sagemaker.Session()

sagemaker_session_bucket=None
if sagemaker_session_bucket is None and sess is not None:

    sagemaker_session_bucket = sess.default_bucket()

try:
    role = sagemaker.get_execution_role()
except ValueError:
    iam = boto3.client('iam')
    role = iam.get_role(RoleName='SageMakerTest')['Role']['Arn']

sess = sagemaker.Session(default_bucket=sagemaker_session_bucket)

print(f"sagemaker role arn: {role}")
print(f"sagemaker bucket: {sess.default_bucket()}")
print(f"sagemaker session region: {sess.boto_region_name}")

```

WARNING:sagemaker:Couldn't call 'get_role' to get Role ARN from role name collab to get Role path.

```

sagemaker role arn: arn:aws:iam::381491942612:role/SageMakerTest
sagemaker bucket: sagemaker-ap-southeast-2-381491942612
sagemaker session region: ap-southeast-2

```

```
In [ ]: ▶ from datasets import load_dataset

# Convert dataset to OAI messages
system_message = """You are an text to SQL query translator. Users will ask you questions in English and you will generate a SQL query based on the provided SCHEMA.
SCHEMA:
{schema}"""

def create_conversation(sample):
    return {
        "messages": [
            {"role": "system", "content": system_message.format(schema=sample["schema"])},
            {"role": "user", "content": sample["question"]},
            {"role": "assistant", "content": sample["answer"]}
        ]
    }

# Load dataset from the hub
dataset = load_dataset("b-mc2/sql-create-context", split="train")
dataset = dataset.shuffle().select(range(12500))

# Convert dataset to OAI messages
dataset = dataset.map(create_conversation, remove_columns=dataset.features)
# split dataset into 10,000 training samples and 2,500 test samples
dataset = dataset.train_test_split(test_size=2500/12500)

print(dataset["train"][345]["messages"])
```

Downloading readme: 0%| | 0.00/4.43k [00:00<?, ?B/s]

Downloading data: 0%| | 0.00/21.8M [00:00<?, ?B/s]

Generating train split: 0 examples [00:00, ? examples/s]

Map: 0%| | 0/12500 [00:00<?, ? examples/s]

```
[{'content': 'You are an text to SQL query translator. Users will ask you questions in English and you will generate a SQL query based on the provided SCHEMA.\nSCHEMA:\nCREATE TABLE table_12962773_4 (no INTEGER, player VARCHAR)', 'role': 'system'}, {'content': 'What No is the player Zoran Erceg', 'role': 'user'}, {'content': 'SELECT MIN(no) FROM table_12962773_4 WHERE player = "Zoran Erceg"', 'role': 'assistant'}]
```

In []: ▶

```

training_input_path = f's3://{sess.default_bucket()}/datasets/text-to-sql'

dataset["train"].to_json(f"{training_input_path}/train_dataset.json", orient='records')
dataset["test"].to_json(f"{training_input_path}/test_dataset.json", orient='records')

print(f"Training data uploaded to:")
print(f"{training_input_path}/train_dataset.json")
print(f"https://s3.console.aws.amazon.com/s3/buckets/{sess.default_bucket()}")

```

Creating json from Arrow format: 0% | 0/10 [00:00<?, ?ba/s]

Creating json from Arrow format: 0% | 0/3 [00:00<?, ?ba/s]

Training data uploaded to:

s3://sagemaker-ap-southeast-2-381491942612/datasets/text-to-sql/train_dataset.json

<https://s3.console.aws.amazon.com/s3/buckets/sagemaker-ap-southeast-2-381491942612/?region=ap-southeast-2&prefix=datasets/text-to-sql/> (<https://s3.console.aws.amazon.com/s3/buckets/sagemaker-ap-southeast-2-381491942612/?region=ap-southeast-2&prefix=datasets/text-to-sql/>)

In []: ▶

hyperparameters, which are passed into the training job

```

hyperparameters = {
    ### SCRIPT PARAMETERS ###
    'dataset_path': '/opt/ml/input/data/training/train_dataset.json', # path to training data
    'model_id': "codellama/Codellama-7b-hf",
    'max_seq_len': 3072,
    'use_qlora': True,
    ### TRAINING PARAMETERS ###
    'num_train_epochs': 3,
    'per_device_train_batch_size': 1,
    'gradient_accumulation_steps': 4,
    'gradient_checkpointing': True,
    'optim': "adamw_torch_fused",
    'logging_steps': 10,
    'save_strategy': "epoch",
    'learning_rate': 2e-4,
    'bf16': False,
    'tf32': True,
    'max_grad_norm': 0.3,
    'warmup_ratio': 0.03,
    'lr_scheduler_type': "constant",
    'report_to': "tensorboard",
    'output_dir': '/tmp/tun',
    'merge_adapters': True,
}

```

```
In [ ]: ▶ from sagemaker.huggingface import HuggingFace

# define Training Job Name
job_name = f'codellama-7b-hf-text-to-sql-exp1'

# create the Estimator
huggingface_estimator = HuggingFace(
    entry_point          = 'run_sft.py',      # train script
    source_dir           = '/content/', # 'https://github.com/philschmid/LLM
    instance_type        = 'ml.t3.medium',    # instances type used for the
    instance_count       = 1,                 # the number of instances us
    max_run              = 2*24*60*60,        # maximum runtime in seconds
    base_job_name        = job_name,          # the name of the training j
    role                 = role,              # Iam role used in training
    volume_size          = 300,              # the size of the EBS volume
    transformers_version = '4.36',           # the transformers version u
    pytorch_version      = '2.1',           # the pytorch_version versio
    py_version           = 'py310',         # the python version used in
    hyperparameters      = hyperparameters, # the hyperparameters passed
    disable_output_compression = True,       # not compress output to sav
    environment           = {
        "HUGGINGFACE_HUB_CACHE": "/tmp/.cache", # set
        "HF_TOKEN": "hf_jwOUrKKKVGoOUrnBQZZQgJTPHyfYTH
    },
)
```

```
In [ ]: ▶ # define a data input dictionary with our uploaded s3 uris
data = {'training': training_input_path}

# starting the train job with our uploaded datasets as input
huggingface_estimator.fit(data, wait=True)
```

```
In [ ]: ▶ from sagemaker.huggingface import get_huggingface_llm_image_uri

# retrieve the Llm image uri
llm_image = get_huggingface_llm_image_uri(
    "huggingface",
    version="1.4.0",
    session=sess,
)

# print ecr image uri
print(f"llm image uri: {llm_image}")
```

llm image uri: 763104351884.dkr.ecr.ap-southeast-2.amazonaws.com/huggingface-pytorch-tgi-inference:2.1.1-tgi1.4.0-gpu-py310-cu121-ubuntu20.04

```
In [ ]: ▶ import json
from sagemaker.huggingface import HuggingFaceModel

# s3 path where the model will be uploaded
# if you try to deploy the model to a different time add the s3 path here
model_s3_path = huggingface_estimator.model_data["S3DataSource"]["S3Uri"]

# sagemaker config
instance_type = "ml.g5.2xlarge"
number_of_gpu = 1
health_check_timeout = 300

# Define Model and Endpoint configuration parameter
config = {
    'HF_MODEL_ID': "/opt/ml/model", # path to where sagemaker stores the model
    'SM_NUM_GPUS': json.dumps(number_of_gpu), # Number of GPU used per replica
    'MAX_INPUT_LENGTH': json.dumps(1024), # Max Length of input text
    'MAX_TOTAL_TOKENS': json.dumps(2048), # Max Length of the generation (inference)
}

# create HuggingFaceModel with the image uri
llm_model = HuggingFaceModel(
    role=role,
    image_uri=llm_image,
    model_data={'S3DataSource':{'S3Uri': model_s3_path, 'S3DataType': 'S3Prefix'},
    env=config
)
```

```
In [ ]: ▶ # Deploy model to an endpoint
# https://sagemaker.readthedocs.io/en/stable/api/inference/model.html#sagemaker.HuggingFaceModel.deploy
llm = llm_model.deploy(
    initial_instance_count=1,
    instance_type=instance_type,
    container_startup_health_check_timeout=health_check_timeout, # 10 minutes
)
```



```
In [ ]: ▶ from transformers import AutoTokenizer
from sagemaker.s3 import S3Downloader

# Load the tokenizer
tokenizer = AutoTokenizer.from_pretrained("codellama/CodeLlama-7b-hf")

# Load the test dataset from s3
S3Downloader.download(f"{training_input_path}/test_dataset.json", ".")
test_dataset = load_dataset("json", data_files="test_dataset.json", split='
random_sample = test_dataset[345]

def request(sample):
    prompt = tokenizer.apply_chat_template(sample, tokenize=False, add_generation_prompt=True)
    outputs = llm.predict({
        "inputs": prompt,
        "parameters": {
            "max_new_tokens": 512,
            "do_sample": False,
            "return_full_text": False,
            "stop": ["<|im_end|>"],
        }
    })
    return {"role": "assistant", "content": outputs[0]["generated_text"].split(prompt)[-1]}

print(random_sample["messages"][1])
request(random_sample["messages"][:2])
```

```

In [ ]: ▶ from transformers import AutoTokenizer
from sagemaker.s3 import S3Downloader

# Load the tokenizer
tokenizer = AutoTokenizer.from_pretrained("codellama/CodeLlama-7b-hf")

# Load the test dataset from s3
S3Downloader.download(f"{training_input_path}/test_dataset.json", ".")
test_dataset = load_dataset("json", data_files="test_dataset.json", split="train")
random_sample = test_dataset[345]

def request(sample):
    prompt = tokenizer.apply_chat_template(sample, tokenize=False, add_generation_prompt=True)
    outputs = llm.predict({
        "inputs": prompt,
        "parameters": {
            "max_new_tokens": 512,
            "do_sample": False,
            "return_full_text": False,
            "stop": ["<|im_end|>"],
        }
    })
    return {"role": "assistant", "content": outputs[0]["generated_text"].split(prompt)[-1]}

print(random_sample["messages"][1])
request(random_sample["messages"][:2])

```

```

In [ ]: ▶ from tqdm import tqdm

def evaluate(sample):
    predicted_answer = request(sample["messages"][:2])
    if predicted_answer["content"] == sample["messages"][2]["content"]:
        return 1
    else:
        return 0

success_rate = []
number_of_eval_samples = 1000
# iterate over eval dataset and predict
for s in tqdm(test_dataset.shuffle().select(range(number_of_eval_samples))):
    success_rate.append(evaluate(s))

# compute accuracy
accuracy = sum(success_rate)/len(success_rate)

print(f"Accuracy: {accuracy*100:.2f}%")

```

```

In [ ]: ▶ llm.delete_model()
llm.delete_endpoint()

```

